

Video Transcript

Video 6: Creating Input for a Hadoop Container (02:13)

To get some experience with Hadoop, we're going to be running a word frequency, a word counting program. We're going to start off by logging into the container, generating some input data, later downloading a program, and after that we're going to be running all of that. So, we'll start off here by accessing the container and creating some of the input. There are a few ways in which you can access that container, but in this case, we're going to take the easy way. We're going to use the graphical user interface, and I'm going to go ahead and select here 'CLI' for command line interface, and I'll go ahead and click it. And as you'll see, it'll open up a new window.

Now that we're inside of it, let's go ahead and list our files. That's just simply a good practice. Then we're going to go ahead and make a directory, and the directory will be input, a directory for input. We're then going to echo into a file inside of input called f1. It was the best of times. It was the worst of times. Go ahead and do that. And now we have that file inside of input. I'm now going to create a second file as you can see there, f2. And if we were to move inside of input, you'd see that we have those two files as expected.

Next, we want to be able to move that content into HDFS, that is the Hadoop file system. I'm going to start off by creating a directory, as you can see there, called input inside of Hadoop. Next, to move our input into the data nodes on HDFS, we're going to run the following, which is basically going to copy our input into HDFS. And so at this point, to review, we have created some input files. We've then taken those input files, and pushed them into the file system, Hadoop's file system. So, we're now ready to take another step.