

Video Transcript

Video 7: Setting Up a Hadoop Container to Run a Word Count Program (02:08)

To count the word frequency on our input, we're going to need a program. Now, we could write that program or we could download it. And this is one of those things that is often done when you stand up one of these platforms and so as you can imagine, there are many of them that have been written. Now, you could download it onto your local machine, your host, and then upload it onto your container. However, in our case, we're going to download it directly from within the container.

We're going to be running curl. As you can see there, the -L option is one you might have not seen before, and this one simply follows the link in case there are any redirects from that initial file that you point to. Then as you can see there, we have an option as well to output and some file name that we're going to give it. So, the address that we're going to be targeting there is one of a repository for Java type programs, you can see there, the long link in blue in that the name of that file is 'hadoop-mapreduce-examples,' a number, and then .jar. And we're going to be saving it as that name as well.

So, let's go ahead and move into the container and download it. So, here we are inside of the container, and I'm going to go ahead and paste that instruction and you can see there that it's the same one that I had on the slide. So, I'm going to go ahead and hit return and you can see there some of the feedback of the data. Now, if I go ahead and list the files, you'll see there that we have there the long file name and that file is now in our in our file system. So, if I list it, you can see there that's perhaps easier to see in that listing. So, go ahead and do the same.

Download that file into your container as mentioned before. If you prefer, you can download it to your file system and then from there push it into the container.